



Linear Regression

DCS310

Sun Yat-sen University

Outline


- Introduction
- Single Feature Case
- Multiple Features Case
- Numerical Optimization

Introduction

- What is regression?

Based on the given features, predict the values of interested variables

- Example: House price prediction



The diagram shows a blue bracket labeled "Features" spanning the first four columns of the table. A red arrow labeled "Interested variable" points to the fifth column of the table.

Size (feet) x_1	# bedrooms x_2	# floors x_3	# years (Ages) x_4	Price (\$ 1000) y
2104	5	1	45	460
1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178
....

Features: 1) size, 2) # bedrooms, 3) # floors, 4) # years

- Mathematically, regression aims at building a function $f(\cdot)$ to model the relation between input data x and supervised value y

$$\hat{y} = f(x_1, x_2, x_3, x_4)$$

- Linear** regression

Restricting the function $f(\cdot)$ to be of **linear form**, *i.e.*,

$$f(x_1, x_2 \cdots x_m) = w_0 + w_1 x_1 + w_2 x_2 + \cdots + w_m x_m$$

- w_k : model parameters
- m : number of features

- Objective

Find a set of parameters $\{w_k\}_{k=1}^m$ so that the prediction

$$\hat{y} = f(x_1, x_2, \dots, x_m)$$

is **as close as possible** to the true y values *for all data samples* in the training dataset

	Size (feet) x_1	# bedrooms x_2	# floors x_3	# years (Ages) x_4	Price (\$ 1000) y
Sample 1 →	2104	5	1	45	460
Sample 2 →	1416	3	2	40	232
Sample 3 →	1534	3	2	30	315
Sample 4 →	852	2	1	36	178

Linear model is very understandable, or saying explainable.

Outline


- Introduction
- **Single Feature Case**
- Multiple Features Case
- Numerical Optimization

Model

- For simplicity, first consider only one feature, e.g., *house size*

Size (feet) x	Price (\$ 1000) y
2104	460
1416	232
852	178
....


- Plot the (x, y) pairs on a plane



- The prediction function reduces to

$$f(x) = w_0 + w_1 x$$

- For different w_0 and w_1 , the function $f(x)$ represents different straight lines




- The objective is to find an appropriate w_0 and w_1 so that **the line is as fit as possible** with the true y 's of all given x

Cost / Loss Function

- Mathematically, the objective can be described as minimizing the cost (loss) function, a.k.a. mean squared error (MSE)

$$L(w_0, w_1) = \frac{1}{n} \sum_{i=1}^n (f(x^{(i)}) - y^{(i)})^2$$


where $x^{(i)}$ and $y^{(i)}$ denote the feature and target value of the i -th training sample; n is the number of training samples



- Substituting $f(x^{(i)}) = w_0 + w_1 x^{(i)}$ into $L(w_0, w_1)$ gives

$$L(w_0, w_1) = \frac{1}{n} \sum_{i=1}^n (w_0 + w_1 x^{(i)} - y^{(i)})^2$$

Remark: To better understand this cost function, we simplify it by setting $w_0 = 0$




- Then, the cost function becomes

$$L(w_1) = \overline{x^2} w_1^2 - 2\overline{xy} w_1 + \overline{y^2}$$


where $\overline{x^2} = \frac{\sum_{i=1}^n (x^{(i)})^2}{n}$, $\overline{xy} = \frac{\sum_{i=1}^n x^{(i)} y^{(i)}}{n}$ and $\overline{y^2} = \frac{\sum_{i=1}^n (y^{(i)})^2}{n}$

- The cost function is a **quadratic function** *w.r.t.* w_1



$$L(w_1) = \overline{x^2}w_1^2 - 2\overline{xy}w_1 + \overline{y^2}$$

- If w_0 is taken into account, the cost function $L(w_0, w_1)$ is still a quadratic function, but is **two-dimensional**



$$L(w_0, w_1) = \frac{1}{n} \sum_{i=1}^n (w_0 + w_1 x^{(i)} - y^{(i)})^2$$

- The best w_0 and w_1 can be found by setting the derivatives to zero

$$\frac{\partial L}{\partial w_0} = \frac{2}{n} \sum_{i=1}^n (w_0 + w_1 x^{(i)} - y^{(i)}) = 0$$

$$\frac{\partial L}{\partial w_1} = \frac{2}{n} \sum_{i=1}^n (w_0 + w_1 x^{(i)} - y^{(i)}) x^{(i)} = 0$$

- It can be derived that the best w_0 and w_1 are

$$w_0 = \frac{\overline{xy\bar{x}} - \overline{x^2}\bar{y}}{\bar{x}^2 - \overline{x^2}}$$

$$w_1 = \frac{\bar{x}\bar{y} - \overline{xy}}{\bar{x}^2 - \overline{x^2}}$$

Outline

- Introduction
- Single Feature Case
- **Multiple Features Case**
- Numerical Optimization

- Training data from single feature to multiple feature case, a.k.a. multivariate linear regression

Size (feet)	Price (\$ 1000)
x	y
2104	460
1416	232
1534	315
852	178
....



Size (feet)	# bedrooms	# floors	# years (Ages)	Price (\$ 1000)
x_1	x_2	x_3	x_4	y
2104	5	1	45	460
1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178
....

- The function of a general linear regression is

$$f(x_1, x_2 \cdots x_m) = w_0 + w_1 x_1 + w_2 x_2 + \cdots + w_m x_m$$

- x_i is the i -th feature
- Working with the scalar form is cumbersome. Reformulating it into a matrix-form gives

$$\underline{f(\mathbf{x}) = \mathbf{x}\mathbf{w}}$$

- $\mathbf{x} = [1, x_1, x_2, \cdots, x_m]$ $\in \mathbb{R}^{1 \times (m+1)}$ is the feature row vector
- $\mathbf{w} = [w_0, w_1, w_2, \cdots, w_m]^T$ $\in \mathbb{R}^{(m+1) \times 1}$ is the parameter column vector

*By setting the first element in \mathbf{x} to be 1, w_0 can be treated **the same** as the other parameters w_k*

Cost Function

- The objective is still to find a \mathbf{w} such that the prediction

$$f(\mathbf{x}^{(i)}) = \mathbf{x}^{(i)} \mathbf{w}$$

is close to the true value $y^{(i)}$, where $\mathbf{x}^{(i)}$ and $y^{(i)}$ is the feature vector and target value of the i -th training sample

Size (feet)	# bedrooms	# floors	# years (Ages)	Price (\$ 1000)
x_1	x_2	x_3	x_4	y
2104	5	1	45	460
1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178
....

- Thus, the cost function can be represented as

$$L(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n (\mathbf{x}^{(i)} \mathbf{w} - y^{(i)})^2$$

- The cost function can be further written as

$$L(\mathbf{w}) = \frac{1}{n} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2$$

where \mathbf{X} and \mathbf{y} are the feature matrix the target vector, defined as

$$\mathbf{X} \triangleq \begin{bmatrix} \mathbf{x}^{(1)} \\ \vdots \\ \mathbf{x}^{(n)} \end{bmatrix} \in \mathbb{R}^{n \times (m+1)}, \quad \mathbf{y} \triangleq \begin{bmatrix} y^{(1)} \\ \vdots \\ y^{(n)} \end{bmatrix} \in \mathbb{R}^{n \times 1}$$

n行 m+1列

w_0

	Size (feet) x_1	# bedrooms x_2	# floors x_3	# years (Ages) x_4	Price (\$ 1000) y
x_0					
1	2104	5	1	45	460
1	1416	3	2	40	232
1	1534	3	2	30	315
1	852	2	1	36	178
1

$\mathbf{X} =$ $\mathbf{y} =$

- The gradient of the cost function w.r.t. \mathbf{w} is

$$\frac{\partial L(\mathbf{w})}{\partial \mathbf{w}} = \frac{2}{n} \mathbf{X}^T (\mathbf{X}\mathbf{w} - \mathbf{y})$$

- Since $L(\mathbf{w})$ is a convex function, its optima can be found by setting

$$\frac{\partial L(\mathbf{w})}{\partial \mathbf{w}} = \frac{2}{n} \mathbf{X}^T (\mathbf{X}\mathbf{w} - \mathbf{y}) = 0$$

特征冗余

- Solving the equation gives

$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

特征数量 m 大于样本数量 n , $\mathbf{X}^T \mathbf{X}$ 不满秩, 则不可逆

- It can be verified that when the number of feature is 1, the result reduces to

$$w_0 = \frac{\overline{xy\bar{x}} - \overline{x^2\bar{y}}}{\bar{x}^2 - \overline{x^2}}, \quad w_1 = \frac{\bar{x}\bar{y} - \overline{xy}}{\bar{x}^2 - \overline{x^2}}$$

- What about $\mathbf{X}^T \mathbf{X}$ is NOT invertible, i.e. not full-rank? Resulting multiple solutions, which one to select? Regularization


Geometric Interpretation

- From the requirement of $X^T(Xw^* - y) = 0$, we can see that

$$\underline{y - Xw^* \perp \text{span}(X)} \quad \text{正交}$$

Xw^* 为 y 在 $\text{span}(X)$ 中的投影

- The result suggests that Xw^* can be understood as *the projection of y onto the space spanned by X*



Outline

- Introduction
- Single Feature Case
- Multiple Features Case
- Numerical Optimization


Gradient Descent

- Analytical solutions do *not always exist*, or evaluating the analytical expression is *computational expensive*
- Then, we should resort to numerical methods, e.g., the gradient descent

梯度下降法

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - r \cdot \left. \frac{\partial L(\mathbf{w})}{\partial \mathbf{w}} \right|_{\mathbf{w}=\mathbf{w}^{(t)}}$$

- r : the learning rate





- Let us take **the single-feature case** and set $w_0 = 0$ as an example, in which the loss function becomes


$$L(w_1) = \frac{1}{n} \sum_{i=1}^n (w_1 x^{(i)} - y^{(i)})^2$$

- The parameter w_1 can be updated as

$$w_1^{(t+1)} = w_1^{(t)} - r \cdot \left. \frac{\partial L(w_1)}{\partial w_1} \right|_{w_1 = w_1^{(t)}}$$




- With **appropriate learning rate**, the model parameter is iteratively updated, and will eventually converge to the optima




- As approaching the optimal value, the gradient becomes smaller and smaller. Thus, *even if the learning rate is fixed*, the updating intervals also **approach 0** as the iteration proceeds, as long as the rate is set appropriately

- If the learning rate is too small, the convergence speed will be very slow




- If the learning rate is too large, the iteration may diverge
- So, setting appropriate learning rate is important




- Now consider the case with both w_0 and w_1

$$w_0^{(t+1)} = w_0^{(t)} - r \cdot \frac{\partial L(w_0, w_1)}{\partial w_0} \Big|_{w_0=w_0^{(t)}, w_1=w_1^{(t)}}$$


$$w_1^{(t+1)} = w_1^{(t)} - r \cdot \frac{\partial L(w_0, w_1)}{\partial w_1} \Big|_{w_0=w_0^{(t)}, w_1=w_1^{(t)}}$$




The function $f(x) = w_0^{(t)} + w_1^{(t)}x$




The contours of $L(w_0, w_1)$ and the track of $((w_0^{(t)}, w_1^{(t)}))$




The function $f(x) = w_0^{(t)} + w_1^{(t)}x$




The contours of $L(w_0, w_1)$ and the track of $((w_0^{(t)}, w_1^{(t)}))$




The function $f(x) = w_0^{(t)} + w_1^{(t)}x$




The contours of $L(w_0, w_1)$ and the track of $((w_0^{(t)}, w_1^{(t)}))$




The function $f(x) = w_0^{(t)} + w_1^{(t)}x$




The contours of $L(w_0, w_1)$ and the track of $(w_0^{(t)}, w_1^{(t)})$




The function $f(x) = w_0^{(t)} + w_1^{(t)}x$




The contours of $L(w_0, w_1)$ and the track of $(w_0^{(t)}, w_1^{(t)})$




The function $f(x) = w_0^{(t)} + w_1^{(t)}x$




The contours of $L(w_0, w_1)$ and the track of $(w_0^{(t)}, w_1^{(t)})$




The function $f(x) = w_0^{(t)} + w_1^{(t)}x$




The contours of $L(w_0, w_1)$ and the track of $(w_0^{(t)}, w_1^{(t)})$




The function $f(x) = w_0^{(t)} + w_1^{(t)}x$




The contours of $L(w_0, w_1)$ and the track of $(w_0^{(t)}, w_1^{(t)})$



The function $f(x) = w_0^{(t)} + w_1^{(t)}x$



The contours of $L(w_0, w_1)$ and the track of $(w_0^{(t)}, w_1^{(t)})$



Stochastic Gradient Descent

- The GD algorithm need to evaluate the gradient of loss w.r.t. model parameters \mathbf{w} *at every iteration*
- Generally, the gradient takes the form

$$\frac{\partial L(\mathbf{w})}{\partial \mathbf{w}} = \frac{1}{n} \sum_{i=1}^n \frac{\partial \ell(\mathbf{w}, \mathbf{x}^{(i)}, y^{(i)})}{\partial \mathbf{w}}$$

- Every iteration requires computing the gradient for *all data samples in the training dataset*

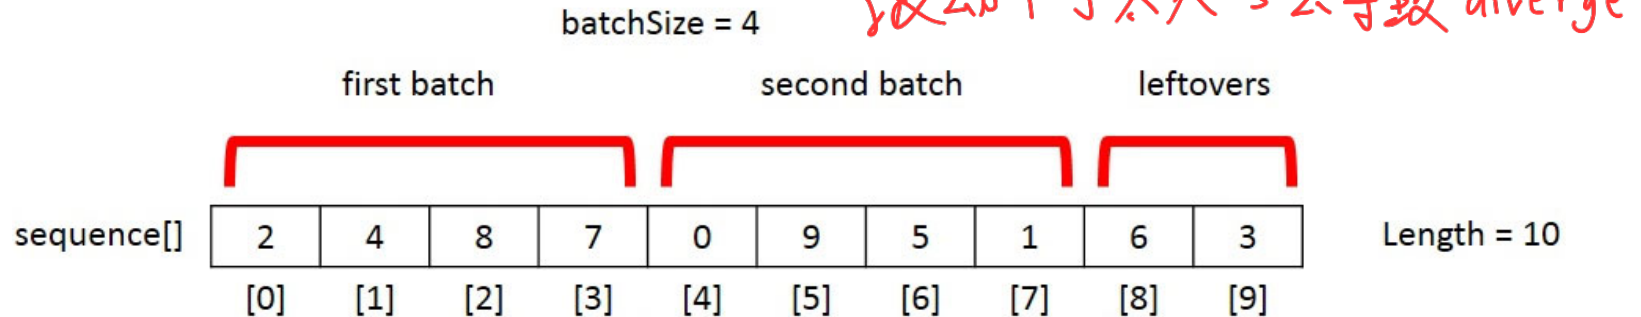
The complexity would be extremely high for large datasets

- To reduce the complexity, we can estimate the gradient $\frac{\partial L(\mathbf{w})}{\partial \mathbf{w}}$ using a small portion of the dataset, *i.e. mini-batch*

- How to obtain the mini-batches?

- Reshuffling

- Segmenting



- Update:

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} + r \cdot \frac{1}{|\mathcal{B}_t|} \sum_{i \in \mathcal{B}_t} \frac{\partial \ell(\mathbf{w}, \mathbf{x}^{(i)}, y^{(i)})}{\partial \mathbf{w}}$$



A noisy estimate to the true gradient

where \mathcal{B}_t is a mini-batch of the dataset at the t -th iteration

Other Optimization Methods

- There also exist many other optimization methods

1) Newton method 牛顿法



Advantages

- No need to manually choose the learning rate
- Faster convergence rate

Disadvantages

- More expensive

- 2) Quasi-Newton methods
- 3) Conjugate gradient method
- 4) Coordinated descent method
- ⋮

These methods generally converge faster than the gradient method, but are more computationally expensive