



Support Vector Machines

DCS310


Sun Yat-sen University

Outline

- Decision Boundaries of Linear Classifiers
- Maximum-Margin Classifier
- Soft Maximum-Margin Classifier
- Support Vector Machine
- Relation to Logistic Regression

Decision Boundaries in Linear Classifiers

- In linear classifiers, the decision boundary is always a hyperplane. The goal is to find the hyperplane that can separate different types of samples




- Logistic regression
 - The decision-boundary hyperplane is found by minimizing the cross-entropy loss

$$L(\mathbf{w}, b) = -y \log(\sigma(\mathbf{w}^T \mathbf{x} + b)) - (1 - y) \log(1 - \sigma(\mathbf{w}^T \mathbf{x} + b))$$

- With the optimal \mathbf{w}^* and b^* , the hyperplane is composed of \mathbf{x} in

$$\{\mathbf{x} | \mathbf{w}^{*T} \mathbf{x} + b^* = 0\}$$



- 1) The hyperplane is *perpendicular* to the vector \mathbf{w}^*
- 2) The distance from the original point to the hyperplane is $-\frac{b^*}{\|\mathbf{w}^*\|}$


- Ideal classifier

➤ The hyperplane is determined by minimizing the loss


$$L(\mathbf{w}, b) = \sum_{\ell=1}^N e\left(y^{(\ell)}(\mathbf{w}^T \mathbf{x}^{(\ell)} + b)\right)$$

$L(\mathbf{w}, b)$ represents the number of misclassified samples

- $y \in \{-1, 1\}$
- $e(z)$ is a jump function, i.e., $e(z) = 0$ if $z \geq 0$; $e(z) = 1$ otherwise




- If the samples are linearly separable, there will be numerous ideal classifiers, which are determined by \mathbf{w}^* and b^*
- Every \mathbf{w}^* and b^* corresponds to a hyperplane




All the hyperplanes can have the loss reduced to zero

Which Hyperplane is the Best?

- The hyperplane in logistic regression is optimal from the perspective of *minimizing the cross-entropy loss*



- All the hyperplanes in the ideal classifier are optimal from the perspective of *minimizing the number of misclassified samples*




Outline

- Decision Boundaries of Linear Classifiers
- **Maximum-Margin Classifier**
- Soft Maximum-Margin Classifier
- Support Vector Machine
- Relation to Logistic Regression

The Maximum-Margin Objective

- To perform well on unseen data, the intuition is to **find a hyperplane that makes the margin as large as possible**



Thanks to the large margin, *it can be expected that an unseen sample is more likely to be classified correctly under such a decision boundary*

How to Represent the Margin?

- The distance from the sample x to the hyperplane \mathcal{H}

- Every sample x can be decomposed as

$$x = m_1 + m_2$$

- m_1 is on the \mathcal{H} , i.e., $w^T m_1 + b = 0$
- $m_2 \perp \mathcal{H}$ and $m_2 \parallel w$


- Thus, we have

$$\begin{aligned} h(x) &\triangleq w^T x + b = w^T (m_1 + m_2) + b \\ &= w^T m_2 \end{aligned}$$

- Due to $m_2 \parallel w$, we can write

$$m_2 = \gamma \cdot \frac{w}{\|w\|},$$

with $|\gamma|$ representing the length of m_2




- Substituting $\mathbf{m}_2 = \gamma \cdot \frac{\mathbf{w}}{\|\mathbf{w}\|}$ into $h(\mathbf{x}) = \mathbf{w}^T \mathbf{m}_2$ gives

$$h(\mathbf{x}) = \gamma \cdot \frac{\mathbf{w}^T \mathbf{w}}{\|\mathbf{w}\|} \quad \Rightarrow \quad \gamma = \frac{h(\mathbf{x})}{\|\mathbf{w}\|}$$

- The distance of a sample on the other side of the hyperplane (i.e. of negative class, $h(\mathbf{x}) < 0$) is

$$\gamma = -\frac{h(\mathbf{x})}{\|\mathbf{w}\|}$$




- The distance of a sample (x, y) to the hyperplane is given by

$$\gamma = \frac{y \cdot h(\mathbf{x})}{\|\mathbf{w}\|} = \frac{y \cdot (\mathbf{w}^T \mathbf{x} + b)}{\|\mathbf{w}\|}$$

where $y \in \{-1, 1\}$

- The margin of a hyperplane under a dataset is given by the minimum distance, *i.e.*,

$$\text{Margin} = \min_{\ell} \frac{y^{(\ell)} \cdot (\mathbf{w}^T \mathbf{x}^{(\ell)} + b)}{\|\mathbf{w}\|}$$



- Thus, the maximum-margin classifier is to find the \mathbf{w}^* and b^* that maximize the margin, *i.e.*,

$$\mathbf{w}^*, b^* = \arg \max_{\mathbf{w}, b} \left\{ \frac{1}{\|\mathbf{w}\|} \min_{\ell} [y^{(\ell)} \cdot (\mathbf{w}^T \mathbf{x}^{(\ell)} + b)] \right\}$$

But how to optimize is unknown

The Transformed Objective Function

- Optimizing another objective function that shares the same optima as the original problem
 - If \mathbf{w}^* and b^* is the optima to $\frac{1}{\|\mathbf{w}\|} \min_{\ell} [y^{(\ell)} \cdot (\mathbf{w}^T \mathbf{x}^{(\ell)} + b)]$, then $\kappa \mathbf{w}^*$ and κb^* is also the optima for all $\kappa \in \mathbb{R}$
 - There must exist an optima $\kappa \mathbf{w}^*$ and κb^* that satisfy the constraints

$$y^{(\ell)} \cdot (\kappa \mathbf{w}^{*T} \mathbf{x}^{(\ell)} + \kappa b^*) \geq 1 \quad \text{for all } \ell = 1, 2, \dots, n$$

Among all \mathbf{w}, b that satisfy $y^{(\ell)} \cdot (\mathbf{w}^T \mathbf{x}^{(\ell)} + b) \geq 1$ for all ℓ , *the \mathbf{w} and b with the smallest $\|\mathbf{w}\|^2$ must be the optima \mathbf{w}^* and b^* that maximizes*

$$\frac{1}{\|\mathbf{w}\|} \min_{\ell} [y^{(\ell)} (\mathbf{w}^T \mathbf{x}^{(\ell)} + b)]$$

- Therefore, the maximum-margin hyperplane can be found by solving the optimization problem below

$$\begin{aligned} \min_{\mathbf{w}, b} \quad & \frac{1}{2} \|\mathbf{w}\|^2 \\ \text{s.t.} \quad & y^{(\ell)} \cdot (\mathbf{w}^T \mathbf{x}^{(\ell)} + b) \geq 1, \quad \text{for } \ell = 1, 2, \dots, N \end{aligned}$$

- This is a convex quadratic optimization problem. Its optimal solution can be found by *numerical methods* efficiently
- With the optimal \mathbf{w}^* and b^* , an unseen data \mathbf{x} can be classified as

$$\hat{y}(\mathbf{x}) = \text{sign}(\mathbf{w}^{*T} \mathbf{x} + b^*)$$

The Equivalent Dual Formulation

- Every convex optimization problem corresponds to an equivalent dual formulation

All contents in this section are extracted from the subject of **convex optimization**

- The **Lagrangian function** of the original optimization problem

$$\mathcal{L}(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{\ell=1}^N a_{\ell} (y^{(\ell)} (\mathbf{w}^T \mathbf{x}^{(\ell)} + b) - 1),$$

where the Lagrange multiplier a_{ℓ} is required to satisfy $a_{\ell} \geq 0$

- The **Lagrange dual function**

$$g(\mathbf{a}) = \min_{\mathbf{w}, b} \mathcal{L}(\mathbf{w}, b, \mathbf{a})$$

- The **dual formulation** of the original optimization problem

$$\begin{array}{ll} \max_{\mathbf{a}} & g(\mathbf{a}) \\ \text{s.t.} & \mathbf{a} \geq \mathbf{0} \end{array}$$

- Deriving the close-form expression (a.k.a. analytical solution) of function $g(\mathbf{a})$, i.e. removing minimization w.r.t. \mathbf{w} and b

- Setting the gradient $\frac{\partial \mathcal{L}}{\partial \mathbf{w}} = \mathbf{0}$ and $\frac{\partial \mathcal{L}}{\partial b} = 0$, which gives (classroom work)

$$\frac{\partial \left(\frac{1}{2} \|\mathbf{w}\|^2 - \sum_{\ell=1}^N a_{\ell} (y^{(\ell)} (\mathbf{w}^T \mathbf{x}^{(\ell)} + b) - 1) \right)}{\partial \mathbf{w}} = \mathbf{0} \Rightarrow \mathbf{w} = \sum_{\ell=1}^N a_{\ell} y^{(\ell)} \mathbf{x}^{(\ell)}$$

$$\frac{\partial \left(\frac{1}{2} \|\mathbf{w}\|^2 - \sum_{\ell=1}^N a_{\ell} (y^{(\ell)} (\mathbf{w}^T \mathbf{x}^{(\ell)} + b) - 1) \right)}{\partial b} = 0 \Rightarrow \sum_{\ell=1}^N a_{\ell} y^{(\ell)} = 0$$

- Substituting them into $\mathcal{L}(\mathbf{w}, b, \mathbf{a})$ gives $g(\mathbf{a}) = \min_{\mathbf{w}, b} \mathcal{L}(\mathbf{w}, b, \mathbf{a})$ as (classroom work)

$$g(\mathbf{a}) = \sum_{\ell=1}^N a_{\ell} - \frac{1}{2} \sum_{\ell=1}^N \sum_{j=1}^N a_{\ell} a_j y^{(\ell)} y^{(j)} \mathbf{x}^{(\ell)T} \mathbf{x}^{(j)}$$

- Then, the dual optimization becomes

$$\begin{array}{ll} \max_{\mathbf{a}} & g(\mathbf{a}) \\ \text{s.t.} & \mathbf{a} \geq \mathbf{0} \text{ and } \sum_{\ell=1}^N a_{\ell} \mathbf{y}^{(\ell)} = \mathbf{0} \end{array}$$

where $g(\mathbf{a}) = \sum_{\ell=1}^N a_{\ell} - \frac{1}{2} \sum_{\ell=1}^N \sum_{j=1}^N a_{\ell} a_j \mathbf{y}^{(\ell)} \mathbf{y}^{(j)T} \mathbf{x}^{(\ell)} \mathbf{x}^{(j)}$

Attention: $g(\mathbf{a})$ no longer contains minimization w.r.t. \mathbf{w} and b

It is a quadratic optimization, and can be solved by *numerical methods*

- Relation between optima \mathbf{w}^* , b^* and optima \mathbf{a}^*

- With the optima \mathbf{a}^* , according to $\mathbf{w} = \sum_{\ell=1}^N a_{\ell} y^{(\ell)} \mathbf{x}^{(\ell)}$, the optimal \mathbf{w}^* is equal to

$$\mathbf{w}^* = \sum_{\ell=1}^N a_{\ell}^* y^{(\ell)} \mathbf{x}^{(\ell)}$$

- Due to $y^{(i)}(\mathbf{w}^{*T} \mathbf{x}^{(i)} + b^*) = 1$ for all samples $(\mathbf{x}^{(i)}, y^{(i)})$ that are on the margin (i.e. $i \in \mathcal{S}$), we can derive that

For each $i \in \mathcal{S}$

$$b^* = \begin{cases} 1 - \mathbf{w}^{*T} \mathbf{x}^{(i)} & \text{if } y^{(i)} = 1 \\ -1 - \mathbf{w}^{*T} \mathbf{x}^{(i)} & \text{if } y^{(i)} = -1 \end{cases}$$

$$= y^{(i)} - \sum_{\ell=1}^N a_{\ell}^* y^{(\ell)} \mathbf{x}^{(\ell)T} \mathbf{x}^{(i)}$$

Mean solution is
more robust

$$b^* = \frac{1}{|\mathcal{S}|} \sum_{i \in \mathcal{S}} \left(y^{(i)} - \sum_{\ell=1}^N a_{\ell}^* y^{(\ell)} \mathbf{x}^{(\ell)T} \mathbf{x}^{(i)} \right)$$

- Maximum-margin classifiers

- Primal version

$$\hat{y}(\mathbf{x}) = \text{sign}(\mathbf{w}^{*T} \mathbf{x} + b^*)$$

- Dual version

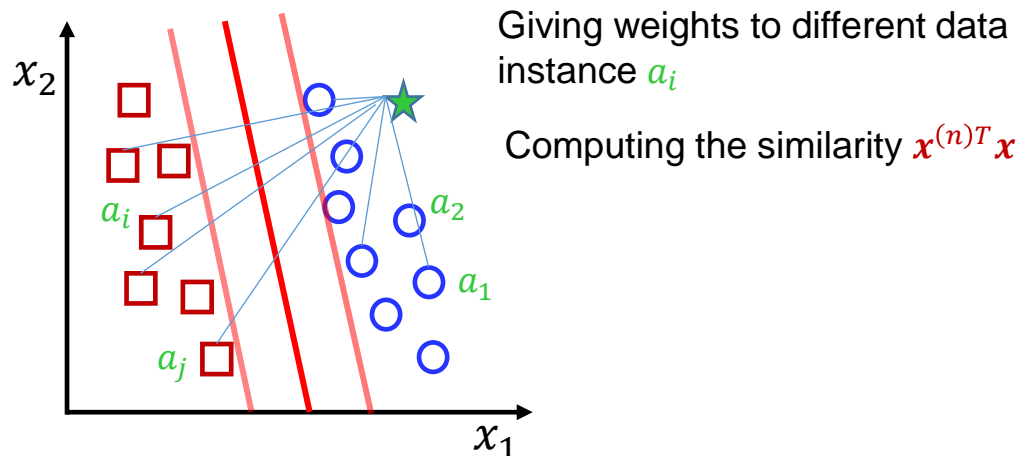
Substituting $\mathbf{w}^* = \sum_{n=1}^N a_n^* y^{(n)} \mathbf{x}^{(n)}$ into the primal version gives

$$\hat{y}(\mathbf{x}) = \text{sign}\left(\sum_{n=1}^N a_n^* y^{(n)} \mathbf{x}^{(n)T} \mathbf{x} + b^*\right)$$

The two classifiers are *equivalent*

$$\hat{y}(\mathbf{x}) = \text{sign} \left(\sum_{n=1}^N \left(a_n^* (\mathbf{x}^{(n)T} \mathbf{x}) \right) \cdot y^{(n)} + b^* \right)$$

- How to understand the dual maximum-margin classifier?
 - For a test \mathbf{x} , computing its similarity with all the training samples $\mathbf{x}^{(n)}$ for $n = 1, \dots, N$ by $\mathbf{x}^{(n)T} \mathbf{x}$
 - Summing all the labels $y^{(n)}$ weighted by the sample similarity $\mathbf{x}^{(n)T} \mathbf{x}$ and the multiplier a_n^*



Comparisons on the Primal and Dual Results

- Optimization (training) complexity

Primal

$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2$$

$$s. t. : y^{(\ell)} \cdot (\mathbf{w}^T \mathbf{x}^{(\ell)} + b) \geq 1, \\ \text{for } \ell = 1, 2, \dots, N$$

of parameter to optimize:
dimension of features

Dual

$$\max_{\mathbf{a}} g(\mathbf{a})$$

$$s. t. : \mathbf{a} \geq \mathbf{0}$$

$$\sum_{\ell=1}^N a_{\ell} y^{(\ell)} = 0$$

of parameter to optimize:
of training samples

In *high-dimensional feature case*, solving the dual problem is more efficient

- Testing complexity

Primal

$$\hat{y}(\mathbf{x}) = \text{sign}(\mathbf{w}^{*T} \mathbf{x} + b^*)$$

Just need **one** inner-product $\mathbf{w}^{*T} \mathbf{x}$

Dual

$$\hat{y}(\mathbf{x}) = \text{sign}\left(\sum_{n=1}^N a_n^* y^{(n)} \mathbf{x}^{(n)T} \mathbf{x} + b^*\right)$$

Need **N** inner-products $\mathbf{x}^{(n)T} \mathbf{x}$
for $n = 1, 2, \dots, N$

At the first glance, the dual classifier looks much more expensive than the primal one

- But it can be shown that **most of a_n^* are 0**

Sparsity in the Lagrange Multiplier a^*

- For any convex optimization problem, the optima satisfies the *KKT conditions*, which, for our problem, are

$$a_n^* \geq 0$$

$$y^{(n)}(\mathbf{w}^{*T} \mathbf{x}^{(n)} + b^*) - 1 \geq 0$$

$$a_n^* [y^{(n)}(\mathbf{w}^{*T} \mathbf{x}^{(n)} + b^*) - 1] = 0$$

The first two conditions come from the original primal and dual problems


- From the last condition, we can see that $a_n^* \neq 0$ only when $y^{(n)}(\mathbf{w}^{*T} \mathbf{x}^{(n)} + b^*) = 1$
- If $\mathbf{x}^{(n)}$ satisfies $y^{(n)}(\mathbf{w}^{*T} \mathbf{x}^{(n)} + b^*) = 1$, it means that it lies on the margin

This kind of samples are called *support vectors*

- Thus, when we classify an unseen sample x as

$$\hat{y}(x) = \text{sign} \left(\sum_{n \in \mathcal{S}} a_n^* y^{(n)} \mathbf{x}^{(n)T} \mathbf{x} + b^* \right),$$

we only need to evaluate the similarity $\mathbf{x}^{(n)T} \mathbf{x}$ between x and the support vectors (samples)




Outline

- Decision Boundaries of Linear Classifiers
- Maximum-Margin Classifier
- **Soft Maximum-Margin Classifier**
- Support Vector Machine
- Relation to Logistic Regression

Non-separable Classes

- The assumption used in the previous maximum-margin classifier

The training samples are linearly separable!!!



- However, *such a hyperplane may not exist*. That is, there is no feasible solution to the optimization problem

$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2$$


$$\text{s.t.}: y^{(n)} \cdot (\mathbf{w}^T \mathbf{x}^{(n)} + b) \geq 1, \quad \text{for } n = 1, 2, \dots, N$$

Soft Maximum Margin

- To address the issue, instead of requiring $y^{(n)} \cdot (\mathbf{w}^T \mathbf{x}^{(n)} + b) \geq 1$ for all $n = 1, \dots, N$, we only require

$$y^{(n)} \cdot (\mathbf{w}^T \mathbf{x}^{(n)} + b) \geq 1 - \xi_n$$

where ξ_n is *slack variable* and $\xi_n \geq 0$
松弛变量



- The objective is not just to minimize $\frac{1}{2} \|\mathbf{w}\|^2$, but also need to minimize the sum of ξ_n , which leads to the objective

$$\frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n$$

where C is used to control the relative importance

- The optimization problem now becomes

$$\min_{\mathbf{w}, b, \xi} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n$$

$$s.t.: y^{(n)} \cdot (\mathbf{w}^T \mathbf{x}^{(n)} + b) \geq 1 - \xi_n,$$

$$\xi_n \geq 0, \quad \text{for } n = 1, 2, \dots, N$$

- Using the same method as before (refer to slide 17), the *dual formulation* can be derived as (homework)

$$\max_{\mathbf{a}} g(\mathbf{a})$$

$$s.t.: a_n \geq 0, \quad a_n \leq C$$

$$\sum_{n=1}^N a_n y^{(n)} = 0$$

$$\text{where } g(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n a_m y^{(n)} y^{(m)} \mathbf{x}^{(n)T} \mathbf{x}^{(m)}$$

- With the optima \mathbf{w}^* and b^* , a sample \mathbf{x} is classified as


$$\hat{y}(\mathbf{x}) = \text{sign}(\mathbf{w}^{*T} \mathbf{x} + b^*)$$

- With the optima \mathbf{a}^* , a sample \mathbf{x} is classified as

$$\hat{y}(\mathbf{x}) = \text{sign}\left(\sum_{n=1}^N a_n^* y^{(n)} \mathbf{x}^{(n)T} \mathbf{x} + b^*\right)$$

Also, the two classifiers above are *equivalent*

- The optima \mathbf{a}^* is sparse, with only elements *within* the margin being nonzero



Outline

- Decision Boundaries of Linear Classifiers
- Maximum-Margin Classifier
- Soft Maximum-Margin Classifier
- **Support Vector Machine**
- Relation to Logistic Regression

Non-linearization

- The maximum-margin classifiers so far are **linear**
- To non-linearizing the model, we can transform the original feature x to another space via the **basis function**

$$\phi: x \rightarrow \phi(x)$$


- The primal maximum-margin optimization problem becomes

$$\min_{w,b,\xi} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n$$

$$s.t.: y^{(n)} \cdot (\mathbf{w}^T \phi(\mathbf{x}^{(n)}) + b) \geq 1 - \xi_n,$$

$$\xi_n \geq 0, \quad \text{for } n = 1, 2, \dots, N$$

Classifier: $\hat{y}(\mathbf{x}) = \text{sign}(\mathbf{w}^{*T} \boldsymbol{\phi}(\mathbf{x}^{(n)}) + b^*)$



- Intuitively, data is easier to be separated in high-dimensional space
- *To achieve better performance, we prefer the dimension of the transformed feature space $\phi(\mathbf{x}^{(n)})$ to be **as high as possible***
- However, the dimension of basis function $\phi(\mathbf{x})$ **cannot be too high**, otherwise the optimization would be very expensive

$$\min_{\mathbf{w}, b, \xi} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n$$

$$s.t.: y^{(n)} \cdot (\mathbf{w}^T \phi(\mathbf{x}^{(n)}) + b) \geq 1 - \xi_n,$$

$$\xi_n \geq 0, \quad \text{for } n = 1, 2, \dots, N$$

- The problem can be solved via its dual form

$$\begin{aligned} \max_{\mathbf{a}} \quad & g(\mathbf{a}) \\ \text{s.t.} \quad & a_n \geq 0, \quad a_n \leq C \\ & \sum_{n=1}^N a_n y^{(n)} = 0 \end{aligned}$$

where $g(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n a_m y^{(n)} y^{(m)} \boldsymbol{\phi}(\mathbf{x}^{(n)})^T \boldsymbol{\phi}(\mathbf{x}^{(m)})$

Classifier: $\hat{y}(\mathbf{x}) = \text{sign} \left(\sum_{n=1}^N a_n^* y^{(n)} \boldsymbol{\phi}(\mathbf{x}^{(n)})^T \boldsymbol{\phi}(\mathbf{x}) + b^* \right)$

- The dimension of \mathbf{a} is *independent of the dimension of $\boldsymbol{\phi}(\cdot)$* , thus the dual form is able to work in very large feature space $\boldsymbol{\phi}(\cdot)$

The dual formulation requires to evaluate the inner product

$$\phi(x^{(n)})^T \phi(x),$$

which is expensive in high-dimensional case

The issue can be addressed by using the *kernel trick*

Kernel Function

- A kernel function is a two-variable function $k(\mathbf{x}, \mathbf{x}')$ that can be expressed as an inner production of some function $\phi(\cdot)$

$$k(\mathbf{x}, \mathbf{x}') = \phi(\mathbf{x})^T \phi(\mathbf{x}')$$

Obviously, $\mathbf{x}^T \mathbf{x}'$ and $\phi(\mathbf{x})^T \phi(\mathbf{x}')$ are kernel functions

- Mercer Theorem:** If a function $k(\mathbf{x}, \mathbf{x}')$ is symmetric positive definite, i.e., 对称正定


$$\int \int g(\mathbf{x}) k(\mathbf{x}, \mathbf{y}) g(\mathbf{y}) d\mathbf{x} d\mathbf{y} \geq 0, \quad \forall g(\cdot) \in L^2,$$

there must exist a function $\phi(\cdot)$ such that $k(\mathbf{x}, \mathbf{x}') = \phi(\mathbf{x})^T \phi(\mathbf{x}')$

If a function $k(\mathbf{x}, \mathbf{x}')$ satisfies the symmetric positive definite condition, it must be a kernel function

- One of the most widely used kernel is the Gaussian kernel, which takes the form

$$k(\mathbf{x}, \mathbf{x}') = \exp \left\{ -\frac{1}{2\sigma^2} \|\mathbf{x} - \mathbf{x}'\|^2 \right\}$$



- The function $\phi(\cdot)$ of Gaussian kernel has **infinite** dimensions

Kernel Trick

- With the kernel function, the dual maximum-margin classifier can be rewritten as

$$\max_{\mathbf{a}} g(\mathbf{a})$$

$$s.t.: \quad a_n \geq 0, \quad a_n \leq C$$

$$\sum_{n=1}^N a_n y^{(n)} = 0$$

$$\text{where } g(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n a_m y^{(n)} y^{(m)} k(\mathbf{x}^{(n)}, \mathbf{x}^{(m)})$$

- The induced classifier

$$\hat{y}(\mathbf{x}) = \text{sign}\left(\sum_{n=1}^N a_n^* y^{(n)} k(\mathbf{x}^{(n)}, \mathbf{x}) + b^*\right)$$

Kernel trick: replacing the $\phi(\mathbf{x})^T \phi(\mathbf{x}')$ with the kernel function $k(\mathbf{x}, \mathbf{x}')$

- The conclusions can be summarized as
 - If $k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^T \mathbf{x}'$, it is a **linear** maximum-margin classifier
 - If $k(\mathbf{x}, \mathbf{x}') = \boldsymbol{\phi}(\mathbf{x})^T \boldsymbol{\phi}(\mathbf{x}')$, it is a **finite-dimensional** nonlinear maximum-margin classifier based on basis functions
 - If $k(\mathbf{x}, \mathbf{x}') = \exp \left\{ -\frac{1}{2\sigma^2} \|\mathbf{x} - \mathbf{x}'\|^2 \right\}$, it is an **infinite-dimensional** nonlinear maximum-margin classifier

Outline

- Decision Boundaries of Linear Classifiers
- Maximum-Margin Classifier
- Soft Maximum-Margin Classifier
- Support Vector Machine
- Relation to Logistic Regression

- In the logistic regression, we minimize the loss

$$\begin{aligned} L(\mathbf{w}, b) &= - \sum_{n=1}^N \left[y^{(n)} \log \sigma(h^{(n)}) + (1 - y^{(n)}) \log (1 - \sigma(h^{(n)})) \right] + \lambda \|\mathbf{w}\|^2 \\ &= \sum_{n=1}^N \log(1 + \exp(-y^{(n)} h^{(n)})) + \lambda \|\mathbf{w}\|^2 \\ &= \sum_{n=1}^N E_{LR}(y^{(n)} h^{(n)}) + \lambda \|\mathbf{w}\|^2 \end{aligned}$$

where $E_{LR}(z) = \log(1 + \exp(-z))$

- In the ideal classifier, we minimize the loss

$$L(\mathbf{w}, b) = \sum_{n=1}^N E_{Ideal}(y^{(n)} h^{(n)}) + \lambda \|\mathbf{w}\|^2$$

where $E_{Ideal}(z) = 0$ if $z \geq 0$; 1 otherwise

- In the linear maximum-margin classifier, we are equivalently minimizing the loss

$$L(\mathbf{w}, b) = \sum_{n=1}^N E_{\infty}(y^{(n)}h^{(n)} - 1) + \frac{1}{2}\|\mathbf{w}\|^2$$

where $E_{\infty}(z) = 0$ if $z \geq 0$; $+\infty$ otherwise

- In the soft linear maximum-margin classifier, we are equivalently minimizing the loss

$$\begin{aligned} L(\mathbf{w}, b) &= C \sum_{n=1}^N E_{SV}(y^{(n)}h^{(n)}) + \frac{1}{2}\|\mathbf{w}\|^2 \\ &= \sum_{n=1}^N E_{SV}(y^{(n)}h^{(n)}) + \lambda\|\mathbf{w}\|^2 \end{aligned}$$

where $E_{SV}(z) = \max(0, 1 - z)$, which is called the *hinge loss*

- So, we can see that the four classifiers can be formulated under the same framework, with the only difference coming from the chosen error function
- The plot of the three error functions

